

What's happening to the Orangutan?

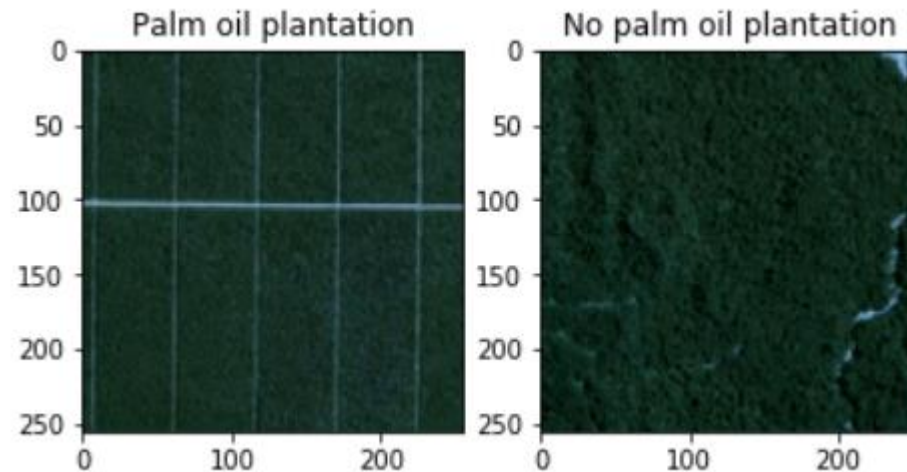
An image classification study
By Cindy Barrientos

The problem

- ▶ Palm oil plantations generate cheap vegetable oils for production
 - ▶ Frozen foods, soap, makeup, detergents
- ▶ Massive deforestation
 - ▶ Orangutans, native to Malaysia and Indonesia, now critically endangered
 - ▶ Orangutan population by year
 - ▶ 288,500 (1973)
 - ▶ 100,000 (2016)
 - ▶ 47,000 (2025, estimated)
- ▶ Satellite imaging to better document location and distribution of world palm oil plantations

Satellite Images

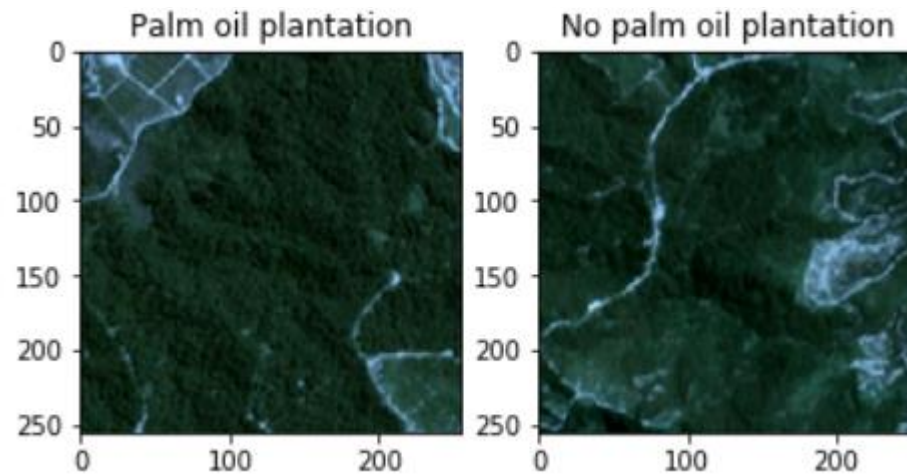
- Satellite imaging to better document location and distribution of world palm oil plantations



Images courtesy of: <https://www.planet.com/>
and <https://www.figure-eight.com/>

Satellite Images

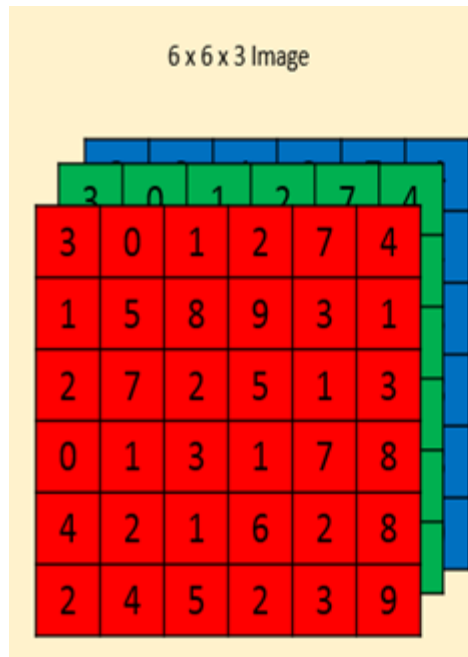
- Satellite imaging to better document location and distribution of world palm oil plantations



Images courtesy of: <https://www.planet.com/>
and <https://www.figure-eight.com/>

Convolutional neural networks

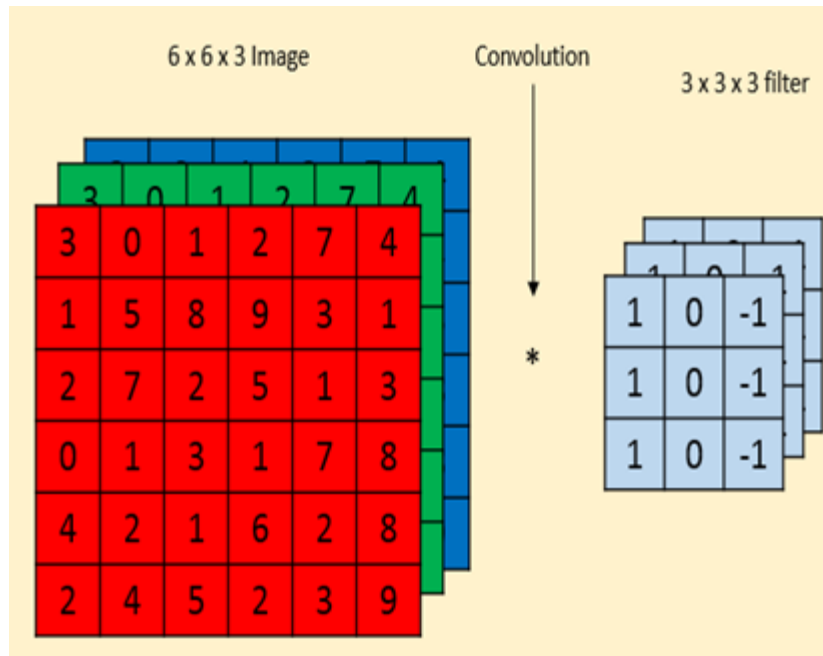
- How an image is perceived by a computer



<https://engmrk.com/convolutional-neural-network-3/>

Convolutional neural networks

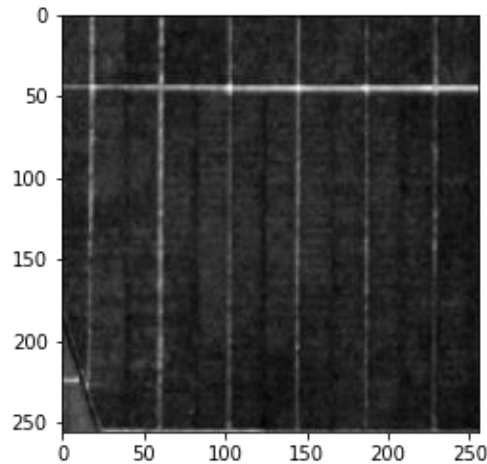
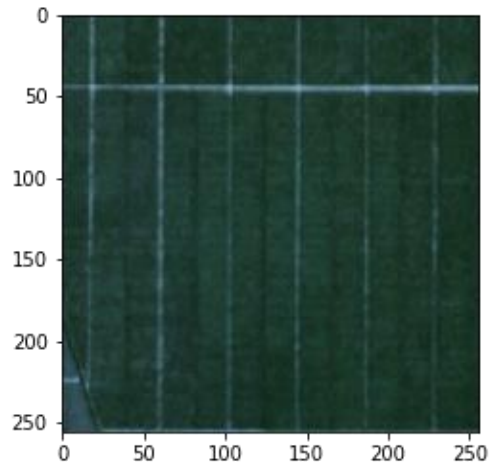
- How an image is perceived by a computer



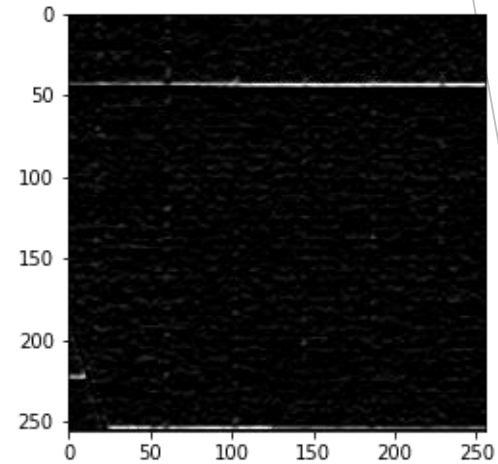
<https://engmrk.com/convolutional-neural-network-3/>

Convolutional neural networks

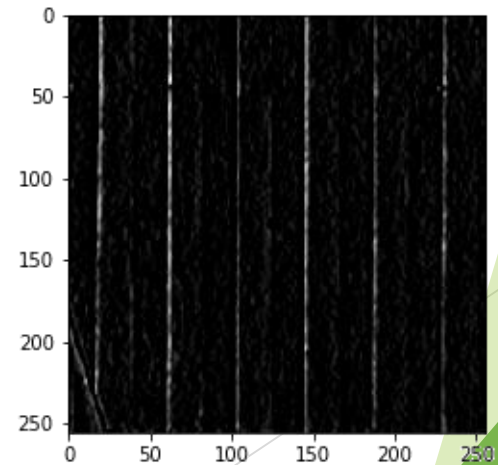
► Edge detection with filters



Grayscale



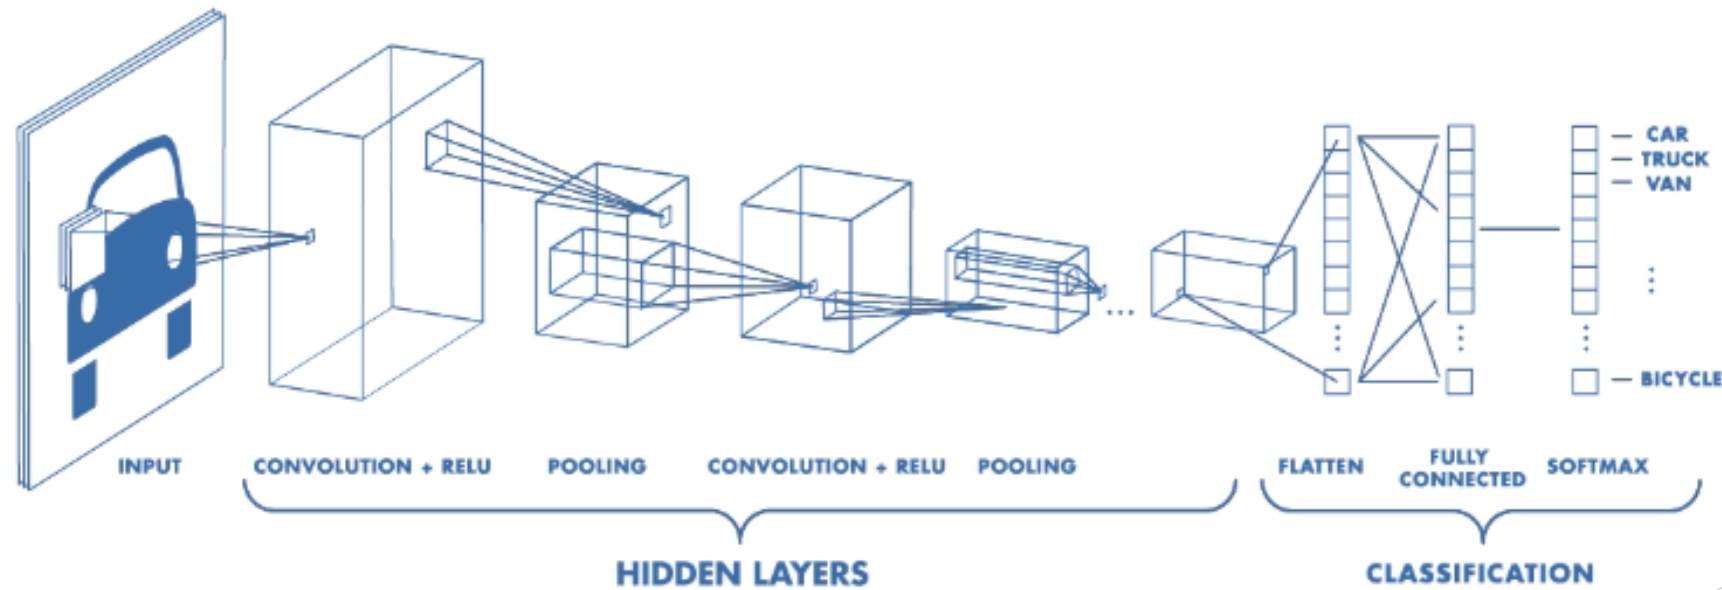
Sobel y



Sobel x

Convolutional neural networks

► Architecture



Class imbalance

```
# list the class sizes  
print("Number of images with palm oil plantations: ", len(os.listdir(has_palm_dir)))  
print("Number of images without palm oil plantations: ", len(os.listdir(no_palm_dir)))
```

```
Number of images with palm oil plantations: 936  
Number of images without palm oil plantations: 14302
```

```
Class distribution of training set:  
(array([0, 1], dtype=int64), array([11450, 741], dtype=int64))
```

```
Class distribution of validation set:  
(array([0, 1], dtype=int64), array([1413, 110], dtype=int64))
```

```
Class distribution of test set:  
(array([0, 1], dtype=int64), array([1439, 85], dtype=int64))
```

Load and transform data

```
# set data transforms
train_transform = transforms.Compose([transforms.Resize(224),
                                     transforms.RandomHorizontalFlip(),
                                     transforms.RandomVerticalFlip(),
                                     transforms.RandomRotation([-30,30]),
                                     transforms.ToTensor(),
                                     transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225))
                                     ])

transform = transforms.Compose([transforms.Resize(224),
                               transforms.ToTensor(),
                               transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225))
                               ])
```

Model: VGG16

- ▶ Winning model designed for the 2014 ImageNet classification challenge
- ▶ Researchers found it to be generalizable to other image sets
- ▶ Uses architecture of small 3x3 convolutional filters
- ▶ Publication: <https://arxiv.org/abs/1409.1556>

Model: VGG16

```
(16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): ReLU(inplace)
(19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace)
  (2): Dropout(p=0.5)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace)
  (5): Dropout(p=0.5)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)
```

Model: VGG16

```
# freeze the features layers
for param in model.features.parameters():
    param.requires_grad = False

classifier = nn.Sequential(OrderedDict([
    ('fc1', nn.Linear(25088, 4096)),
    ('relu1', nn.ReLU()),
    ('dropout1', nn.Dropout(p=0.25)),
    ('fc2', nn.Linear(4096, 4096)),
    ('relu2', nn.ReLU()),
    ('dropout2', nn.Dropout(p=0.25)),
    ('output', nn.Linear(4096, 2))
]))

model.classifier = classifier

print(model)
```

Testing

► Confusion matrix

```
array([[1431,    8],  
       [    8,   77]], dtype=int64)
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	1439
1	0.91	0.91	0.91	85

Class imbalance

- To offset class imbalance, some later trainings were done on 5x duplicated minority class images

Number of images with palm oil plantations: 4680

Number of images without palm oil plantations: 14302

Class distribution of training set:

```
(array([0, 1], dtype=int64), array([11470, 3716], dtype=int64))
```

Class distribution of validation set:

```
(array([0, 1], dtype=int64), array([1441, 457], dtype=int64))
```

Class distribution of test set:

```
(array([0, 1], dtype=int64), array([1391, 507], dtype=int64))
```

Testing

- Confusion matrix of identical model hyperparameters
 - 5 epochs, SGD optimizer, 0.01 learning rate

Original dataset

```
array([[1431,  8],  
       [  8, 77]], dtype=int64)
```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	1439
1	0.91	0.91	0.91	85

5x minority class images

```
array([[1430, 18],  
       [ 17, 433]], dtype=int64)
```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	1435
1	1.00	0.83	0.90	463

Submission predictions

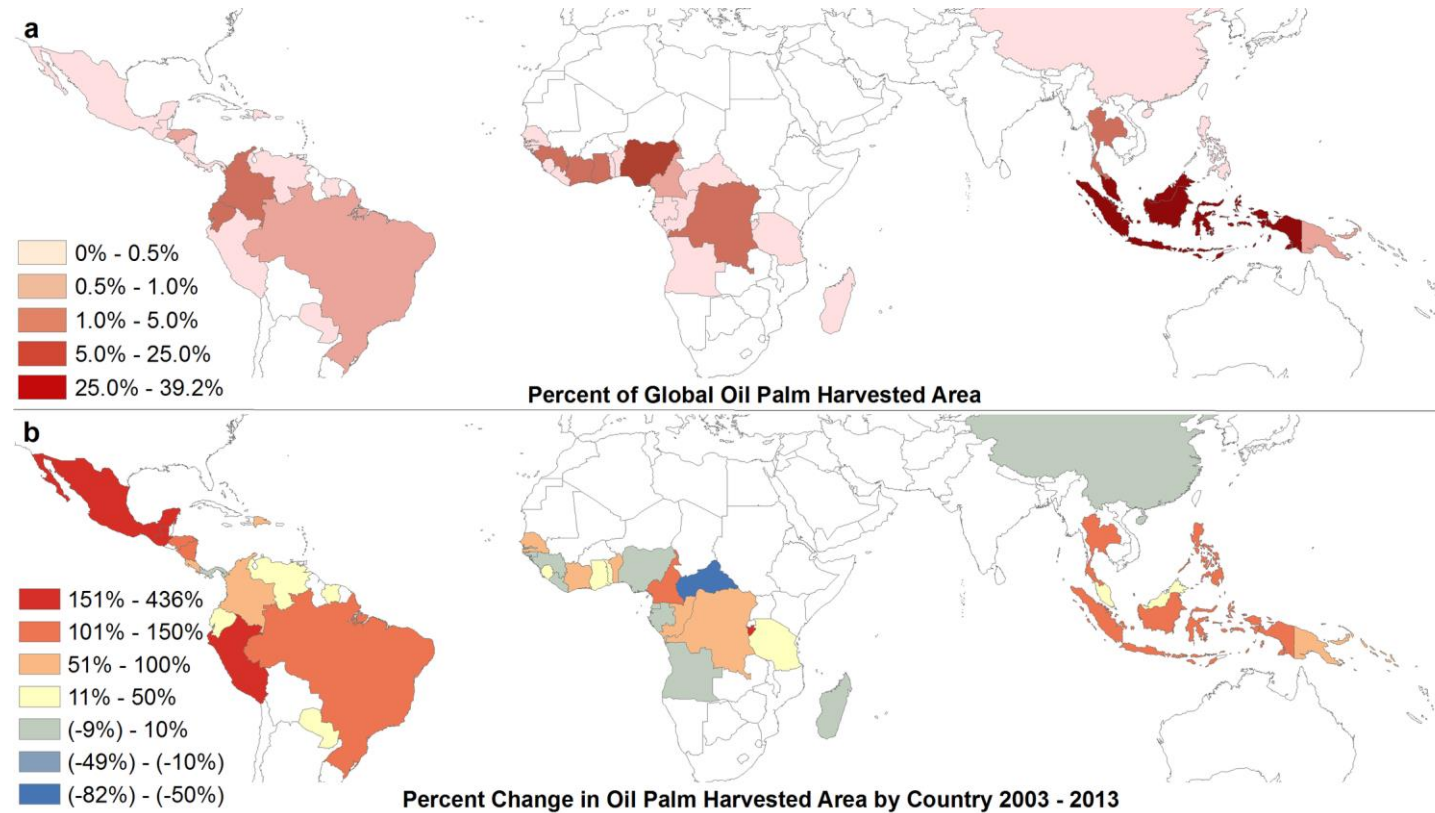
model name (.pt)	Layers + output	Dropout	Optimizer	Learn rate	Epochs	Augmentation	Testval	Leaderboard
vgg16-5ep-sgd01	3	0.5	SGD	0.01	5	Resize(224), Randomhorizontalflip(), randomverticalflip()	0.99081	0.99627
vgg16-5ep-dr75-sgd01	3	0.75	SGD	0.01	5	Resize(224), Randomhorizontalflip(), randomverticalflip()	0.99212	0.59369
vgg16-5ep-dr75-sgd01	3	0.75	SGD	0.01	5	Resize(224), Randomhorizontalflip(), randomverticalflip()	0.98752	0.99652
vgg16-5ep-adam001	3	0.25	Adam	0.001	5	Resize(224), Randomhorizontalflip(), randomverticalflip()	0.99212	0.98439
vgg16bn-5ep-sgd01	3	0.25	SGD	0.01	5	Resize(224), Randomhorizontalflip(), randomverticalflip()	0.99015	0.99469
vgg16-fc4-5ep-sgd01	4	0.25	SGD	0.01	5	Resize(224), Randomhorizontalflip(), randomverticalflip()	0.99146	0.99603
vgg16-imgaug1-5ep-sgd01	3	0.25	SGD	0.01	5	Resize(224), Randomhorizontalflip(), randomverticalflip(), colorjitter(brightness=0.5, contrast=0.75)	0.99212	0.99634
vgg16-5ximg-5ep-sgd01	3	0.25	SGD	0.01	5	Resize(224), Randomhorizontalflip(), randomverticalflip(), randomrotation(-30,30)	0.98156	0.99207

Tweaking for improvements

- ▶ Model
 - ▶ Densenet151, Resnet121
- ▶ Adding fully connected layers
- ▶ Optimizer
 - ▶ SGD vs Adam
 - ▶ Tuning optimal learning rate (perhaps with fast.ai)
- ▶ Class imbalance
- ▶ Image processing

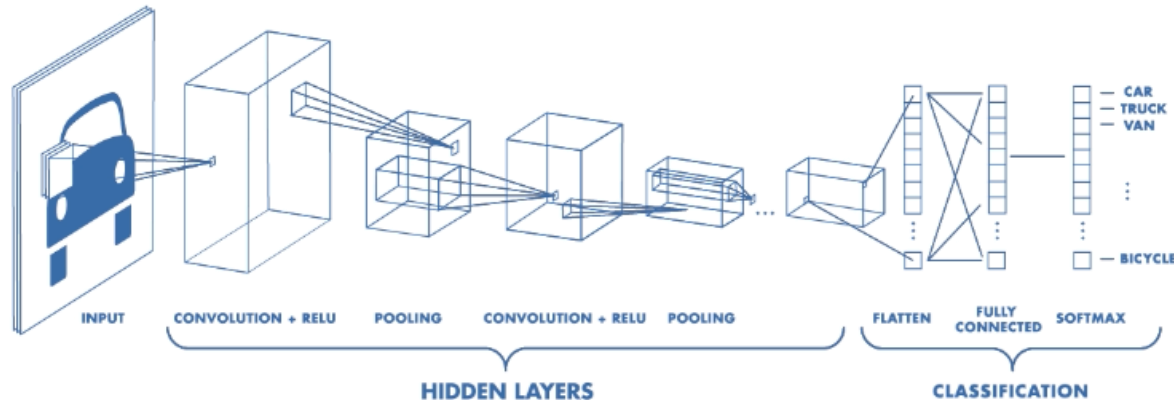
Q&A

Palm oil plantation map



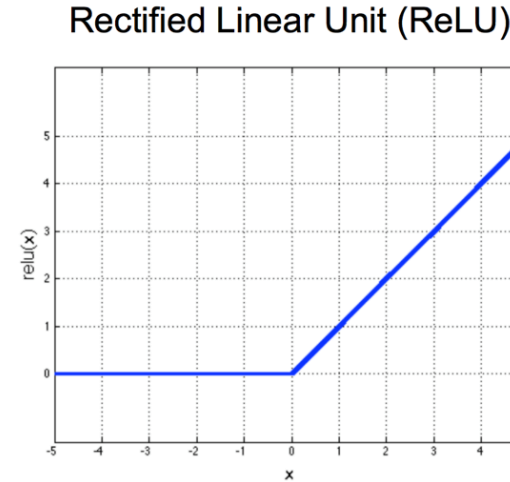
Convolutional neural networks

- ▶ Learning
 - ▶ Loss (mistakes between predicted and actual)
 - ▶ Backpropagation (how a particular weight is contributing to loss)
 - ▶ Optimization (how the algorithm will correct itself)



Activation functions

$$\text{relu}(x) \begin{cases} 0, & \text{if } x < 0 \\ x, & \text{if } x \geq 0 \end{cases}$$



- ▶ Purpose of activation function:
 - ▶ Scales output of layers so they are consistent, small values
 - ▶ Efficiency in model training
- ▶ ReLU activation
 - ▶ Rectified Linear Unit
 - ▶ Clips negative values to zero
 - ▶ Generally should be used after every hidden layer for consistent, positive results

Loss: Cross entropy loss

- ▶ Difference between predicted and actual classification
- ▶ First adds a softmax activation function
 - ▶ Turns scores into probabilities
- ▶ Then applies negative log likelihood loss (NLLLoss)
- ▶ Returns the average loss over an entire batch